

Advanced PostgreSQL & Background Jobs for Lending System

Project Overview

This guide outlines 4 powerful features to implement using PostgreSQL's advanced capabilities and background jobs for a collateral-based lending system. Each feature is designed to provide deep exposure to complex database operations and real-world business logic.

Feature 1: Intelligent Collateral Revaluation Engine

PostgreSQL Features to Explore

Window Functions & Common Table Expressions (CTEs)

- **Moving Averages:** Calculate 7-day, 30-day price trends
- **Price Volatility:** Use LAG/LEAD functions for price change analysis
- **Ranking Functions:** ROW_NUMBER(), RANK(), DENSE_RANK() for price comparisons
- **Recursive CTEs:** Build hierarchical risk scoring models

JSON Operations

- Store market data in JSONB columns for flexible schema
- Query nested JSON structures for price history
- Index JSON fields for performance
- Use JSON aggregation functions for reporting

Triggers & Stored Procedures

- Automatic revaluation triggers on price updates
- Custom functions for complex valuation logic
- Event-driven collateral status updates
- Audit trail generation

Materialized Views

- Pre-computed aggregations for dashboard performance

- Refresh strategies for real-time vs batch updates
- Indexing strategies for materialized views

Background Job Implementation

Daily/Hourly Jobs

Job: Market Price Fetcher

- Frequency: Every hour during market hours
- Tasks:
 - * Fetch live gold/silver prices from APIs
 - * Store in collateral_valuations table
 - * Calculate price deltas and volatility

Job: Portfolio Revaluation

- Frequency: Daily at market close
- Tasks:
 - * Recalculate all collateral current values
 - * Update loan LTV ratios
 - * Identify margin call triggers

Job: Risk Alert Generator

- Frequency: Real-time (event-driven)
- Tasks:
 - * Monitor LTV threshold breaches
 - * Generate automated notifications
 - * Update loan risk scores

Complex SQL Concepts to Master

Moving Averages with Window Functions

-- Example: Calculate 30-day moving average for gold prices

```
SELECT
    valuation_date,
    market_price,
    AVG(market_price) OVER (
        ORDER BY valuation_date
        ROWS BETWEEN 29 PRECEDING AND CURRENT ROW
    ) as moving_avg_30day
FROM collateral_valuations
WHERE collateral_id = 'gold_standard'
ORDER BY valuation_date;
```

Percentile Calculations for Risk Assessment

- PERCENTILE_CONT() for continuous distributions
- PERCENTILE_DISC() for discrete distributions
- Custom risk scoring based on historical volatility

Advanced Indexing Strategies

- Partial indexes for active collateral only
- Composite indexes for time-series queries
- GIN indexes for JSON columns
- Expression indexes for calculated fields

Feature 2: Payment Behavior Analytics & Default Prediction

PostgreSQL Features to Explore

Advanced Aggregations

- **GROUPING SETS**: Multi-dimensional analysis
- **ROLLUP**: Hierarchical aggregations
- **CUBE**: All possible combinations analysis
- **Custom aggregates**: Domain-specific calculations

Statistical Functions

- CORR(): Correlation between variables
- REGR_SLOPE(): Linear regression analysis
- STDDEV(): Standard deviation calculations
- PERCENTILE functions for distribution analysis

Time-Series Analysis

- DATE_TRUNC() for period grouping
- EXTRACT() for temporal components
- AGE() for duration calculations
- Custom date arithmetic functions

Pattern Matching

- Regular expressions for payment pattern detection
- Fuzzy matching for customer identification
- String similarity functions

Background Job Implementation

Weekly/Monthly Jobs

Job: Payment Pattern Analyzer

- Frequency: Weekly
- Tasks:
 - * Cohort analysis of customer payment behavior
 - * Identify early/late payment patterns
 - * Calculate customer reliability scores

Job: Default Risk Calculator

- Frequency: Monthly
- Tasks:
 - * Run predictive models on payment history
 - * Update customer risk classifications
 - * Generate early warning reports

Job: Seasonal Trend Analyzer

- Frequency: Quarterly
- Tasks:
 - * Identify seasonal payment patterns
 - * Adjust risk models for seasonal factors
 - * Optimize collection strategies

Complex Queries You'll Master

Cohort Analysis

-- Example: Monthly cohort analysis of loan performance

```
WITH loan_cohorts AS (
    SELECT
        DATE_TRUNC('month', disbursement_date) as cohort_month,
        id as loan_id,
        borrower_id
    FROM loans
),
payment_periods AS (
    SELECT
        lc.cohort_month,
        lc.loan_id,
        DATE_TRUNC('month', p.payment_date) as payment_month,
        SUM(p.amount) as total_paid
```

```
FROM loan_cohorts lc
JOIN payments p ON lc.loan_id = p.loan_id
GROUP BY 1,2,3
)
SELECT
    cohort_month,
    (payment_month - cohort_month) as period_number,
    COUNT(DISTINCT loan_id) as active_loans,
    AVG(total_paid) as avg_payment
FROM payment_periods
GROUP BY 1,2
ORDER BY 1,2;
```

Customer Segmentation

- RFM analysis (Recency, Frequency, Monetary)
- K-means clustering using PostgreSQL
- Customer lifetime value calculations

Feature 3: Dynamic Portfolio Risk Management

PostgreSQL Features to Explore

Partitioning

- Range partitioning by date for historical data
- Hash partitioning for large customer datasets
- Partition pruning optimization
- Automated partition management

Full-Text Search

- Document analysis for loan applications
- Search optimization with GIN indexes
- Custom text search configurations
- Ranking and relevance scoring

Advanced Constraints

- Custom check constraints for business rules
- Exclusion constraints for overlapping data

- Deferred constraints for complex validations
- Trigger-based constraint enforcement

Background Job Implementation

Real-time/Batch Jobs

Job: Portfolio Concentration Monitor

- Frequency: Real-time (streaming)
- Tasks:
 - * Monitor collateral type concentrations
 - * Alert on excessive exposure limits
 - * Suggest portfolio rebalancing

Job: Stress Test Runner

- Frequency: Daily
- Tasks:
 - * Run Monte Carlo simulations
 - * Test portfolio under adverse scenarios
 - * Generate risk reports for management

Job: Regulatory Compliance Checker

- Frequency: Continuous
- Tasks:
 - * Validate compliance with lending regulations
 - * Generate automated compliance reports
 - * Alert on policy violations

Advanced Concepts

Portfolio Correlation Analysis

-- Example: Correlation between different collateral types

```
SELECT
    c1.type as collateral_type_1,
    c2.type as collateral_type_2,
    CORR(cv1.total_value, cv2.total_value) as price_correlation
FROM collateral_valuations cv1
JOIN collateral_valuations cv2 ON cv1.valuation_date = cv2.valuation_date
JOIN collaterals c1 ON cv1.collateral_id = c1.id
JOIN collaterals c2 ON cv2.collateral_id = c2.id
WHERE c1.type != c2.type
    AND cv1.valuation_date >= CURRENT_DATE - INTERVAL '1 year'
GROUP BY 1,2
HAVING COUNT(*) > 100 -- Sufficient data points
ORDER BY 3 DESC;
```

Risk Scenario Modeling

- Value-at-Risk (VaR) calculations
- Stress testing with recursive queries
- Portfolio optimization algorithms

Feature 4: Automated Loan Lifecycle Management

PostgreSQL Features to Explore

Event Sourcing Patterns

- Immutable event logs
- State reconstruction from events
- Temporal queries and point-in-time analysis
- Event store optimization

State Machines in SQL

- Finite state machine implementation
- Valid transition enforcement
- State history tracking
- Complex workflow orchestration

Queue Systems

- PostgreSQL as message broker
- Job queuing with proper locking
- Priority queue implementation
- Dead letter queue handling

Concurrent Processing

- Advisory locks for job coordination
- MVCC optimization
- Lock contention analysis
- Deadlock prevention strategies

Background Job Implementation

Event-Driven Jobs

Job: Loan Status Manager

- Trigger: State change events
- Tasks:
 - * Validate state transitions
 - * Update related entities
 - * Generate notifications

Job: Payment Processor

- Trigger: Payment received events
- Tasks:
 - * Apply payments to loans
 - * Update schedules
 - * Calculate outstanding amounts

Job: Interest Calculator

- Frequency: Daily
- Tasks:
 - * Compound interest calculations
 - * Update loan balances
 - * Generate interest charges

Job: Overdue Manager

- Frequency: Daily
- Tasks:
 - * Identify overdue loans
 - * Apply penalties
 - * Trigger collection processes

Advanced Database Patterns

Event Sourcing Implementation

-- Example: Event store table design

```
CREATE TABLE loan_events (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  loan_id TEXT NOT NULL,  
  event_type TEXT NOT NULL,  
  event_data JSONB NOT NULL,  
  event_version INTEGER NOT NULL,  
  created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),  
  created_by TEXT NOT NULL  
);
```

-- Rebuild loan state from events

```
WITH loan_events_ordered AS (  
  SELECT * FROM loan_events  
  WHERE loan_id = $1
```



```

        ORDER BY event_version
    )
    SELECT
        loan_id,
        jsonb_object_agg(key, value) as current_state
    FROM (
        SELECT
            loan_id,
            jsonb_each_text(event_data) as (key, value),
            ROW_NUMBER() OVER (PARTITION BY jsonb_each_text(event_data).key ORDER BY event_version) as rn
        FROM loan_events_ordered
    ) ranked_events
    WHERE rn = 1
    GROUP BY loan_id;

```

Saga Pattern for Distributed Transactions

- Long-running transaction management
- Compensation actions for rollbacks
- Process orchestration
- Failure recovery mechanisms

Recommended Learning Progression

Phase 1: Foundation (Weeks 1-2)

Start with Feature 2: Payment Behavior Analytics

Why This First:

- Most straightforward to implement
- Great introduction to window functions
- Immediate business value
- Foundation for other features

Key Learning Objectives:

- Master window functions and CTEs
- Understand statistical analysis in SQL
- Learn background job patterns
- Develop reporting dashboards

Phase 2: Real-Time Processing (Weeks 3-4)

Move to Feature 1: Collateral Revaluation Engine

Building On:

- Analytics knowledge from Phase 1
- Introduction to real-time processing
- Complex trigger development
- Performance optimization

Key Learning Objectives:

- Real-time data processing
- Trigger and stored procedure development
- JSON operations and indexing
- Materialized view optimization

Phase 3: Advanced Workflows (Weeks 5-6)

Tackle Feature 4: Lifecycle Management

Advanced Concepts:

- Most complex business logic
- State machine implementations
- Event-driven architecture
- Concurrent processing patterns

Key Learning Objectives:

- Event sourcing patterns
- State machine design
- Queue-based processing
- Advanced locking strategies

Phase 4: Enterprise Scale (Weeks 7-8)

Master Feature 3: Portfolio Risk Management

Culmination of Learning:

- Combines all previous knowledge

- Highest performance challenges
- Enterprise-level reporting
- Advanced mathematical modeling

Key Learning Objectives:

- Partitioning and scaling strategies
- Complex mathematical operations
- Advanced indexing techniques
- Performance tuning mastery

Technical Stack Mastery

PostgreSQL Advanced Features

Query Optimization

- Execution plan analysis
- Index strategy development
- Query rewriting techniques
- Performance monitoring

Data Types and Operations

- JSONB operations and indexing
- Array operations and functions
- Full-text search capabilities
- Geographic data types (PostGIS)

Concurrency and Reliability

- Transaction isolation levels
- Locking mechanisms
- Connection pooling
- Backup and recovery strategies

Background Job Patterns

Scheduling Strategies

- Cron-based scheduling
- Event-driven triggers
- Priority queue management
- Load balancing techniques

Error Handling

- Retry mechanisms with exponential backoff
- Dead letter queue implementation
- Circuit breaker patterns
- Monitoring and alerting

Performance Optimization

- Batch processing techniques
- Parallel job execution
- Resource management
- Memory optimization

Business Value and Learning Outcomes

Immediate Benefits

- **Risk Reduction:** Automated monitoring and early warning systems
- **Operational Efficiency:** Reduced manual processes and human error
- **Regulatory Compliance:** Automated compliance checking and reporting
- **Customer Experience:** Faster processing and better service

Technical Skills Developed

- **Database Architecture:** Advanced PostgreSQL features and optimization
- **System Design:** Event-driven and real-time processing systems
- **Performance Tuning:** Query optimization and scaling strategies
- **Business Logic:** Complex financial calculations and risk modeling

Career Impact

- **Full-Stack Database Skills:** From basic queries to advanced optimization
- **Real-World Experience:** Production-ready financial system development
- **Problem-Solving:** Complex business logic implementation
- **Architecture:** Scalable system design principles

Getting Started Checklist

Prerequisites

- ☐ PostgreSQL 14+ installed and configured
- ☐ Background job framework (e.g., Sidekiq, Bull, Agenda)
- ☐ Prisma CLI and basic understanding
- ☐ Node.js/Python for job processing

Phase 1 Setup (Payment Analytics)

- ☐ Set up development database with sample data
- ☐ Install job processing framework
- ☐ Create basic analytics queries
- ☐ Implement first background job

Development Best Practices

- ☐ Version control for database migrations
- ☐ Test data generation scripts
- ☐ Performance monitoring setup
- ☐ Documentation and code comments

Success Metrics

- ☐ Query performance benchmarks
- ☐ Job processing throughput
- ☐ System reliability metrics
- ☐ Business value demonstration

Conclusion

This comprehensive guide provides a structured approach to mastering advanced PostgreSQL features and background job processing through real-world financial system development. Each phase builds upon the previous, ensuring a deep understanding of both technical concepts and business applications.

The progression from basic analytics to complex event-driven systems will provide valuable experience applicable to any enterprise-level application development. Focus on understanding the underlying principles rather than just implementing features, as this knowledge will be transferable to many other domains and technologies.

Remember: The goal is not just to build features, but to develop a deep understanding of data-driven system architecture and advanced database operations that will serve you throughout your career.